

MI3.22

Advanced Programming for HPC

Master ICT, USTH, 2nd year

Aveneau Lilian

lilian.aveneau@univ-poitiers.fr
XLIM/SIC/IG, CNRS, Computer Science Department
University of Poitiers

Year 2018/2019

Lecture 3 – Scan & other Patterns

- Scan
- Segmented Scan
- Others Parallel Patterns


Overview

- Scan
- Segmented Scan
- Others Parallel Patterns


SCAN

It consists to apply a binary, associative and commutative operator to all elements of a given input X of size n , to **compute a list of new values**:

$$\text{SCAN}(X) = \{Y_i\}_{i=0}^{n-1} = \left\{ \bigoplus_{k=0}^i X_k \right\}_{i=0}^{n-1} \quad \text{where } \bigoplus_{k=i}^i = X_i$$




?




SCAN

It consists to apply a binary, associative and commutative operator to all elements of a given input X of size n , to **compute a list of new values**:

$$\text{SCAN}(X) = \{Y_i\}_{i=0}^{n-1} = \left\{ \bigoplus_{k=0}^i X_k \right\}_{i=0}^{n-1} \quad \text{where } \bigoplus_{k=i}^i = X_i$$




$$\text{with } [i, j] = X_i \oplus X_{i+1} \oplus \dots \oplus X_j$$



SCAN

It consists to apply a binary, associative and commutative operator to all elements of a given input X of size n , to **compute a list of new values**:


$$\text{SCAN}(X) = \{Y_i\}_{i=0}^{n-1} = \left\{ \bigoplus_{k=0}^i X_k \right\}_{i=0}^{n-1} \quad \text{where } \bigoplus_{k=i}^i = X_i$$



SCAN

It consists to apply a binary, associative and commutative operator to all elements of a given input X of size n , to **compute a list of new values**:


$$\text{SCAN}(X) = \{Y_i\}_{i=0}^{n-1} = \left\{ \bigoplus_{k=0}^i X_k \right\}_{i=0}^{n-1} \quad \text{where } \bigoplus_{k=i}^i = X_i$$



SCAN

It consists to apply a binary, associative and commutative operator to all elements of a given input X of size n , to **compute a list of new values**:

$$\text{SCAN}(X) = \{Y_i\}_{i=0}^{n-1} = \left\{ \bigoplus_{k=0}^i X_k \right\}_{i=0}^{n-1} \quad \text{where } \bigoplus_{k=i}^i = X_i$$



- The swiss-knife of any parallel programmer!
- In thrust: `thrust::inclusive_scan` with different possibilities

PREFIX-SCAN

It consists to apply a binary associative commutative operator to all elements of a given input X of size n , to **compute a list of new values**:

$$\text{PREFIX-SCAN}(X) = \left\{ \bigoplus_{k=0}^{i-1} X_k \right\}_{i=0}^{n-1}$$

where $\forall j < i$, $\bigoplus_i^j = 0$ or any other initialization value

- A *shifted* SCAN, in fact!

Example with $\oplus = +$

X_i	1	1	1	1	1	1	1	1
SCAN	1	2	3	4	5	6	7	8
PRESCAN	0	1	2	3	4	5	6	7

- It is also the swiss-knife of any parallel programmer!
- In thrust: `thrust::exclusive_scan` with different possibilities


Overview

- Scan
- **Segmented Scan**
- Others Parallel Patterns

SEGMENTED SCAN I

It is a SCAN or a PREFIX-SCAN done **per segment** of a given array ...

Example:



where S defines the segments, and so:

- $Y_0 = X_0$, $Y_1 = X_0 \oplus X_1$, and $Y_2 = X_0 \oplus X_1 \oplus X_2$
- $Y_3 = X_3$, $Y_4 = X_3 \oplus X_4$
- $Y_5 = X_5$, $Y_6 = X_5 \oplus X_6$, $Y_7 = X_5 \oplus X_6 \oplus X_7$ and $Y_8 = X_5 \oplus X_6 \oplus X_7 \oplus X_8$

SEGMENTED SCAN II

Obviously, for PREFIX-SEG-SCAN, we have

- $Y_0 = 0$, $Y_1 = X_0$, and $Y_2 = X_0 \oplus X_1$
- $Y_3 = 0$, $Y_4 = X_3$
- $Y_5 = 0$, $Y_6 = X_5$, $Y_7 = X_5 \oplus X_6$ and $Y_8 = X_5 \oplus X_6 \oplus X_7$

Notice that it is a SCAN (or PREFIX-SCAN) using a particular operator:

$$\otimes : \{\mathbb{N} \times T\}^2 \rightarrow T$$

$$\otimes(\{S_i, X_i\}, \{S_j, X_j\}) = \begin{cases} X_i \oplus X_j & \text{if } S_i = S_j, \\ X_j & \text{otherwise.} \end{cases}$$

In practice, computation cost is a little higher:

- Need to read more data (the segment id)
- Avoid it when possible (prefer SCAN)


With Thrust, implemented in `thrust::inclusive_scan_by_key` and `thrust::exclusive_scan_by_key`

Overview

- Scan
- Segmented Scan
- Others Parallel Patterns

SPLIT or PARTITION

Allows to **separate an array into two segments** using a given predicate:



Using Thrust, you can use `thrust::partition` and variants ...

Notice that it can be written using others patterns:


```

1 PARTITION( $X, P$ )
2   {  $X$  and  $P$  are arrays of size  $N$ ,  $X$  being data and  $P_i = \{0, 1\}$  a boolean array }
3    $Y, Z$ : arrays of size  $N$ 
4   FOR each PE  $i$  in parallel
5      $Y_i \leftarrow P_i$ 
6      $Z_i \leftarrow P_i - 1$  { so contains 0 or -1 }
7    $Y \leftarrow \text{PREFIX-SCAN}(Y)$ 
8    $Z \leftarrow \text{REVERSE-SCAN}(Z)$  { scan in reverse order }
9   FOR each PE  $i$  in parallel
10    IF NOT  $P_i$  THEN  $Y_i \leftarrow N + Z_i$ 
11  return SCATTER( $X, Y$ )

```

COMPACT

COMPACT is similar to PARTITION, but forgetting the second part (for which predicate $P_i = 0$) ...



- Here, destination is not completely modified ...
- Result has variable length!

With Thrust, take a look at *stream compaction* at http://thrust.github.io/doc/group__stream__compaction.html

SORT

SORT is a generic pattern for **sorting data**, given an order

- Many algorithms exist, for parallel computers
- An old one: RICHARD COLE's sorting machine 1986, "Parallel Merge Sort"
- Bitonic sort
- Parallel Quicksort
- Probably the most efficient: radix sort (see labwork)

In thrust, you can use `thrust::sort` which implements a fast radix-sort with bitonic sort into blocks to accelerate the whole process ...