

MI3.22

Advanced Programming for HPC

Master ICT, USTH, 2nd year

Aveneau Lilian

lilian.aveneau@univ-poitiers.fr
XLIM/SIC/IG, CNRS, Computer Science Department
University of Poitiers

Year 2019/2020

Lecture 4 – CUDA: Atomics & Stream

- Atomic Operations
- Locking Memory Pages
- CUDA Streams
- Host Memory Access since GPU

Overview

- Atomic Operations
- Locking Memory Pages
- CUDA Streams
- Host Memory Access since GPU

Compute Capability

Previous SON's examples use only basic CUDA capacities

Notion of compute capacity

- Growing versions: 1.0, 1.1, 1.2, 1.3, 2.0, ... 7.5 (Turing)
- “Russian dolls” model: **version 2.0 includes the previous, and so on ...**
- Basically it defines the **instruction set, types, behavior ...**
- List on *CUDA Zone*

http://www.nvidia.fr/object/cuda_home_new_en.html

Examples:	GeForce 470 GTX	2.0	Quadro M2020	5.2
	GTX Titan Black	3.5	GeForce GTX 1080	6.1
	Tesla K80	3. 7	GeForce 2080	7.5

Which capacities for Atomic Operation

- $SM \geq 1.1 \Rightarrow$ Atomic Operations **on global memory**
- $SM \geq 1.2 \Rightarrow$ Atomic Operations **on shared memory**

Compilation: option `-arch sm_11` or `-arch sm_12` or above

Atomic Operations in a nutshell

Parallelism induces new needs

Example with following instruction: `x++`

- Load `x` - Add 1 - Write `x`
⇒ Operation *read-modify-write*
- Race condition using many PEs: execution order problem
- Example with only 2 PEs: $C_6^3 = \binom{6}{3} = 20$ possibilities!
 - R1 - A1 - W1 - R2 - A2 - W2 : *ok*
 - R1 - R2 - A1 - W1 - A2 - W2 : *error*
 - Only two good orders: 90% lead to a bad behavior!

Solutions

- Mutex, Semaphore, ... not CUDA, or using `atomicCas()`
- Synchronization (barrier, ...): *too slow, not always possible (only inside a same block)*
- Atomic Operations, *the best way with CUDA*

Example: histogram computation

Generic tool, with multiple uses

- Image analysis, compression, computer vision, AI learning, audio codecs ...
- Example with "PROGRAMMING WITH CUDA C" :

	A	C	D	G	H	I	M	N	O	P	R	T	U	W
3	2	2	1	2	1	2	2	1	1	1	2	1	1	1

Example of histogram computation, on a random byte array

```
const int SIZE=100*(1<<20); // 100 Mega
int main () {
    // Allocate and random initialization
    unsigned char *buffer =
        big_random_block( SIZE );
    // Histogram computation
    unsigned int histo[256];
    memset ( histo, 0, sizeof(histo) );
    for (int i=0; i<SIZE; i++)
        ++ histo[ buffer[i] ];
}
```

```
// Naive computation check
long histoCount = 0;
for ( int i=0; i<256; i++ )
    histoCount += histo[ i ];
// Computation ok -> histoCount == SIZE ...
std::cout<< "Histogram sum: "
    << histoCount << " == " << SIZE
    << "?" << std::endl;
return 0;
}
```

On my computer, computation time is 55 ms (without initialization)

Histogram on GPU

It is a *read-modify-write* operation! So atomic add:

```
--global-- void computeHisto ( const unsigned char*const buffer ,
                               unsigned int*const histo , const long size ) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;
    while ( i < size ) {      atomicAdd( &histo[buffer[i]], 1 );      i += stride;  }
}

int main () { // initializations , GPU allocations ...
    unsigned char *buffer = big_random_block( SIZE );
    unsigned char*dev_buffer;
    HANDLE_ERROR( cudaMalloc( (void**)&dev_buffer , SIZE ));
    HANDLE_ERROR( cudaMemcpy( dev_buffer , buffer , SIZE , cudaMemcpyHostToDevice )); // upload
    unsigned int*dev_histo;
    HANDLE_ERROR( cudaMalloc( (void**)&dev_histo , 256*sizeof( int ) ));
    HANDLE_ERROR( cudaMemset( dev_histo , 0 , 256*sizeof( int ) )); // Yes, it exists!
    cudaDeviceProp props; // calcul heuristique du nombre de blocs
    HANDLE_ERROR( cudaGetDeviceProperties( &props , 0 ));
    const int blocs = props.multiProcessorCount;
    computeHisto<<<blocs*2,256>>> ( dev_buffer , dev_histo , SIZE );
    // download the histogram , cleanup GPU resources
    unsigned int histo[ 256 ];
    HANDLE_ERROR( cudaMemcpy( histo , dev_histo , 256*sizeof(int) , cudaMemcpyDeviceToHost ));
    HANDLE_ERROR( cudaFree( dev_buffer )); HANDLE_ERROR( cudaFree( dev_histo ));
    long histoCount = 0; // (too) Naive check the result
    for (int i=0; i<256; i++) histoCount += histo[ i ];
    if ( histoCount != SIZE ) { std::cerr << "Erreur de calcul de l'histogramme!!\n"; }
    return 0;
}
```

Computation time: 67 ms on Quadro M2200 (Maxwell device)

We can do better!

Solution to calculate faster: **add more atomic instructions!**

```
--global-- void computeHisto ( const unsigned char*const buffer ,
                               unsigned int*const histo ,
                               const long size )
{
    // shared histogram inside each block
    --shared-- unsigned int temp[ 256 ];
    // initialization , and so barrier
    temp[ threadIdx.x ] = 0;
    __syncthreads();
    // initialization and "classical" computation
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    const int stride = blockDim.x * gridDim.x;
    while ( i < size ) {
        atomicAdd( &temp[buffer[i]], 1 ); // Warning: SM >= 12 !!!
        i += stride;
    }
    // Waiting all block threads before to write the result in GLOBAL memory
    __syncthreads();
    atomicAdd( &histo[threadIdx.x], temp[ threadIdx.x ] ); // Now SM >= 11
}
```

Computation time: **5 ms on Quadro M2200**

- More difference on old Fermi devices
- Kepler/Maxwell improves the atomic operations

Locks (or mutex)

No specific instruction, but can be done (as *spinlock* on CPU):

- Use an integer stored on GPU's DRAM
- AtomicCAS(): **Compare And Store**

```
#ifndef __LOCK_H__
#define __LOCK_H__
#include "common.h"

class Lock {
    int *mutex; // plays the "mutex" role, accessible by all threads
public:
    Lock( void ) { // only for the host (GPU receives a lock by copy)
        HANDLE_ERROR( cudaMalloc( (void**)&mutex, sizeof(int) ) ); // allocated on GPU
        HANDLE_ERROR( cudaMemset( mutex, 0, sizeof(int) ) );
    }

    ~Lock( void ) {
        cudaFree( mutex );
    }

    __device__ void lock( void ) { // How works atomicCAS( ptr, old, new ):
        while( atomicCAS( mutex, 0, 1 ) != 0 ); // IF *ptr == old THEN *ptr=new; returns old;
                                                // ELSE returns *ptr;
    }

    __device__ void unlock( void ) {
        atomicExch( mutex, 0 ); // Atomically set *mutex=0
    }
};
#endif
```

Example with dot product

```
--global__ void dot( Lock lock, float *a, float *b, float *c ) { // receive lock + c
extern __shared__ float cache[]; // allocated at kernel call (triple bracket)
int tid = threadIdx.x + blockIdx.x * blockDim.x;
const int cacheIndex = threadIdx.x;
float temp = 0.f;
while (tid < N) { temp += a[tid] * b[tid]; tid += blockDim.x * gridDim.x; }
cache[cacheIndex] = temp;
__syncthreads(); // threads synchronization into this block
int i = blockDim.x >> 1; // to reduce, blockDim.x must be a power of 2
while (i != 0) {
    if (cacheIndex < i) cache[cacheIndex] += cache[cacheIndex + i];
    __syncthreads();
    i >>= 1;
}
if ( cacheIndex == 0 ) {
    lock.lock(); // take the spinlock (waiting loop)
    *c += cache[0]; // -> we are in an exclusive area, so we can modify c!
    lock.unlock(); // unlock the spinlock
}
}
```

Call:

```
// allocate c: only one float!
HANDLE_ERROR( cudaMemcpy( dev_c, &c, sizeof(float), cudaMemcpyHostToDevice ) );
Lock lock; // The constructor is called on CPU
dot <<<blocksPerGrid, threadsPerBlock, sizeof(float)*threadsPerBlock>>>
    ( lock, dev_a, dev_b, dev_c );
// No more sum on CPU! We can use directly "c" ...
```

While it is not really faster (26,8 ms for 256 Mega), it is now fully on GPU

...

Overview

- Atomic Operations
- Locking Memory Pages
- CUDA Streams
- Host Memory Access since GPU

cudaHostAlloc()

Allocating into the heap: malloc() or cudaHostAlloc()?

Using malloc()

- Classical mechanism: allocates **pages** on main memory
- Memory is uploaded/downloaded on disk (**swap mechanism**)
- Leads to slow-down when **page fault** occurs

Using cudaHostAlloc()

- Allocation of pin-pages, **so with constant address!**
- So, can be accessed directly to/from the GPU **using DMA (Direct Memory Access)**
- No more slow-down (no page fault), but: **system saturation risk**
- Such an allocation provides “**pinned memory**”

In practice memory transfer always **USES** pinned memory: we can use it explicitly to avoid CPU copy to pinned memory ...

Benchmark (1/3)

```
#include <iostream>
#include "common.h"
const unsigned NbLoops = 100u;
float cuda_malloc_test( const long size , const bool up )
{
    ChronoGPU chr;

    char *a = new char[size]; // or using cudaHostAlloc()
    char *dev_a;
    HANDLE_ERROR( cudaMalloc( (void**)&dev_a , size ) );

    chr.start();
    for (int i=0; i<NbLoops; i++) {
        if (up)
            HANDLE_ERROR( cudaMemcpy( dev_a , a , size , cudaMemcpyHostToDevice ) );
        else
            HANDLE_ERROR( cudaMemcpy( a , dev_a , size , cudaMemcpyDeviceToHost ) );
    }
    chr.stop();
    float elapsedTime = chr.elapsedTime();

    delete a;
    HANDLE_ERROR( cudaFree( dev_a ) );

    return elapsedTime;
}
```

Benchmark (2/3)

```
float cuda_host_alloc_test( const long size , const bool up )
{
    ChronoGPU chr;

    char *a, *dev_a;
    // Allocation in "pinned" mode
    HANDLE_ERROR( cudaHostAlloc( (void**)&a, size, cudaHostAllocDefault ) );
    HANDLE_ERROR( cudaMalloc( (void**)&dev_a, size ) );

    chr.start();
    for (int i=0; i<NbLoops; i++) {
        if (up)
            HANDLE_ERROR( cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice ) );
        else
            HANDLE_ERROR( cudaMemcpy( a, dev_a, size, cudaMemcpyDeviceToHost ) );
    }
    chr.stop();
    float elapsedTime = chr.elapsedTime();

    // Take a look at the cleanup using cudaFreeHost() ...
    HANDLE_ERROR( cudaFreeHost( a ) );
    HANDLE_ERROR( cudaFree( dev_a ) );

    return elapsedTime;
}
```

Benchmark (3/3)

```

int main( void ) {
    const long SIZE = 1<<28; // so 256 Mb
    const float MB = static_cast<float>( NbLoops * (SIZE>>20) );

    // Bench using classical 'cudaMalloc', uploading (host to GPU)
    float elapsedTime = cuda_malloc_test( SIZE, true );
    cout<< "Time_using_cudaMalloc:_"<<elapsedTime<<"_ms\n";
    cout<< "\tMB/s_during_copy_up:_"<<(MB/(elapsedTime/1000.f))<<endl;

    // Bench using classical 'cudaMalloc', downloading (GPU to host)
    elapsedTime = cuda_malloc_test( SIZE, false );
    cout<< "Time_using_cudaMalloc:_"<<elapsedTime<<"_ms\n";
    cout<< "\tMB/s_during_copy_down:_"<<(MB/(elapsedTime/1000.f))<<endl;

    // Bench using 'cudaHostAlloc', uploading (host to GPU)
    elapsedTime = cuda_host_alloc_test( SIZE, true );
    cout<< "Time_using_cudaHostAlloc:_"<<elapsedTime<<"_ms\n";
    cout<< "\tMB/s_during_copy_up:_"<<(MB/(elapsedTime/1000.f))<<endl;

    // Bench using 'cudaHostAlloc', downloading (GPU to host)
    elapsedTime = cuda_host_alloc_test( SIZE, false );
    cout<< "Time_using_cudaHostAlloc:_"<<elapsedTime<<"_ms\n";
    cout<< "\tMB/s_during_copy_down:_"<<(MB/(elapsedTime/1000.f))<<endl;

    return 0;
}
//Time using cudaMalloc: 4908.47 ms
//      MB/s during copy up: 5215.47
// Time using cudaMalloc: 5106.2 ms
//      MB/s during copy down: 5013.51
// Time using cudaHostAlloc: 2460.02 ms
//      MB/s during copy up: 10406.4
// Time using cudaHostAlloc: 2596.96 ms
//      MB/s during copy down: 9857.67

```

Overview

- Atomic Operations
- Locking Memory Pages
- **CUDA Streams**
- Host Memory Access since GPU

Introduction

Pinned memory useful, but it is mainly a necessity for Streams ...

Streams

- Queue of operations on GPU, **which are executed in the order ...**
- Contains: kernel, memory transfers, event recording
- Utility: mainly to execute different tasks in parallel
 - Needs `cudaDeviceProperties.deviceOverlap == true`

How to use them

- 1 Dedicated structure `cudaStream_t`
- 2 Asynchronous memory copies `cudaMemcpyAsync()`
- 3 Launch kernel **using** `kernel<<<bl,th,0,stream>>>()`
- 4 Synchronization using `cudaSynchronize(stream)`

Warning

Asynchronous memory: needs **pinned memory**

Example mono-stream (1/3)

```

#include <iostream>
#include "common.h"
#include <random> // Random number generator, with periodicity 2^19937-1
using namespace std;

const int N = 1024*1024 ; // 1 Mo, the number of threads
const int FULL_DATA_SIZE = N * 20; // 20 data per thread

__global__ void kernel( int *a, int *b, int *c ) {
    const int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < N) { // In fact, the following calculation does not matter ;-)
        const int idx1 = (idx + 1) & 0xFF;
        const int idx2 = (idx + 2) & 0xFF;
        const float as = (a[idx] + a[idx1] + a[idx2]) / 3.0f;
        const float bs = (b[idx] + b[idx1] + b[idx2]) / 3.0f;
        c[ idx ] = (as + bs) * .5f;
    }
}

int main( void ) {
    cudaDeviceProp prop; // We verify that GPU allows copy + kernel in ||
    int whichDevice;
    HANDLE_ERROR( cudaGetDevice( &whichDevice ) );
    HANDLE_ERROR( cudaGetDeviceProperties( &prop, whichDevice ) );
    if (!prop.deviceOverlap) {
        cout<< " Device will not handle overlaps, so no speed up from streams\n";
    }

    ChronoGPU      chr;
    float          elapsedTime;

    cudaStream_t stream; // Stream identifier

```

Example mono-stream (2/3)

```
int *host_a , *host_b , *host_c ;
int *dev_a , *dev_b , *dev_c ;
```

```
// Create/initialize the stream
HANDLE_ERROR( cudaStreamCreate( &stream ) );
```

```
// Allocation on GPU
HANDLE_ERROR( cudaMalloc( (void**)&dev_a , N * sizeof(int) ) );
HANDLE_ERROR( cudaMalloc( (void**)&dev_b , N * sizeof(int) ) );
HANDLE_ERROR( cudaMalloc( (void**)&dev_c , N * sizeof(int) ) );
```

```
// Pinned memory allocation , needed by streams
const unsigned int TRUE_SIZE = FULL_DATA_SIZE * sizeof( int );
HANDLE_ERROR( cudaHostAlloc( (void**)&host_a , TRUE_SIZE , cudaHostAllocDefault ) );
HANDLE_ERROR( cudaHostAlloc( (void**)&host_b , TRUE_SIZE , cudaHostAllocDefault ) );
HANDLE_ERROR( cudaHostAlloc( (void**)&host_c , TRUE_SIZE , cudaHostAllocDefault ) );
```

```
// Initialization of some data
auto mt_rand = std::bind( std::uniform_int_distribution<int>(0,FULL_DATA_SIZE) , std::mt1993
    for (int i=0; i<FULL_DATA_SIZE; i++) {
        host_a[i] = mt_rand();
        host_b[i] = mt_rand();
    }
```

```
// Starts the chronometer
chr.start();
```

Example mono-stream (3/3)

```

// Loop on the data, using batch of size N (lot?)
const int N4 = sizeof( int ) * N;
for (int i=0; i<FULL.DATA.SIZE; i+= N) {
    // Asynchronous copies of data from host's RAM to GPU
    HANDLE_ERROR( cudaMemcpyAsync( dev_a, host_a+i, N4, cudaMemcpyHostToDevice, stream ) );
    HANDLE_ERROR( cudaMemcpyAsync( dev_b, host_b+i, N4, cudaMemcpyHostToDevice, stream ) );

    kernel<<<N/256,256,0,stream>>>( dev_a, dev_b, dev_c ); // 0: no shared meory

    // Data copy from GPU to host
    HANDLE_ERROR( cudaMemcpyAsync( host_c+i, dev_c, N4, cudaMemcpyDeviceToHost, stream ) );
}
// Waits the last copy to host's RAM
HANDLE_ERROR( cudaStreamSynchronize( stream ) );

chr.stop();
elapsedTime = chr.elapsedTime();
cout<< "Time_taken: _" << elapsedTime << " _ms\n";

// Spring cleanup
HANDLE_ERROR( cudaFreeHost( host_a ) );
HANDLE_ERROR( cudaFreeHost( host_b ) );
HANDLE_ERROR( cudaFreeHost( host_c ) );
HANDLE_ERROR( cudaFree( dev_a ) );
HANDLE_ERROR( cudaFree( dev_b ) );
HANDLE_ERROR( cudaFree( dev_c ) );
HANDLE_ERROR( cudaStreamDestroy( stream ) ); // Note the release of the stream

return 0;
}

```

Version with two streams

Main loop: launch the two streams ...

```
for (int i=0; i<FULL_DATA_SIZE; i+= N*2) { // half iteration!  
    // Works with the first stream  
    HANDLE_ERROR( cudaMemcpyAsync( dev_a0 , host_a+i, N4, cudaMemcpyHostToDevice, stream0 ));  
    HANDLE_ERROR( cudaMemcpyAsync( dev_b0 , host_b+i, N4, cudaMemcpyHostToDevice, stream0 ));  
    kernel<<< N/256, 256, 0, stream0 >>>( dev_a0 , dev_b0 , dev_c0 );  
    HANDLE_ERROR( cudaMemcpyAsync( host_c+i, dev_c0 , N4, cudaMemcpyDeviceToHost, stream0 ));  
  
    // Works with the second stream  
    HANDLE_ERROR( cudaMemcpyAsync( dev_a1 , host_a+i+N, N4, cudaMemcpyHostToDevice, stream1));  
    HANDLE_ERROR( cudaMemcpyAsync( dev_b1 , host_b+i+N, N4, cudaMemcpyHostToDevice, stream1));  
    kernel<<< N/256, 256, 0, stream1 >>>( dev_a1 , dev_b1 , dev_c1 );  
    HANDLE_ERROR( cudaMemcpyAsync( host_c+i+N, dev_c1 , N4, cudaMemcpyDeviceToHost, stream1));  
}
```

Result

- Same timings!
- In fact, we need to understand the GPU:
 - 2 virtual processors: **one for copies, the other for kernel**
 - Each one has its own **order queue**
 - Waiting to respect the stream order: **induced constraint, we need to take it into account in our stream usage!**
- Solution: **launch 2nd kernel BEFORE downloading the 1st**

Overview

- Atomic Operations
- Locking Memory Pages
- CUDA Streams
- Host Memory Access since GPU

Zero-Copy Host Memory

Principle: pinned memory can be directly accessed from the GPU

Example: dot product

- We reuse the previous code, without allocating memory on GPU
- Instead, we "map" host memory to GPU
 - Function `cudaHostGetDevicePointer(void**, void*, 0)`
- It needs unmovable memory (so, pinned memory)

Replace the classical host allocation

- `cudaHostAllocPortable`: accessible from each threads on CPU
- `cudaHostAllocWriteCombined` :
 - Do not use L1 and L2 cache: slower when reading,
 - But faster when writing – no more page fault handling
- `cudaHostAllocMapped`: Grants access since GPU

Results

On my laptop, 2 to 4 times faster!